

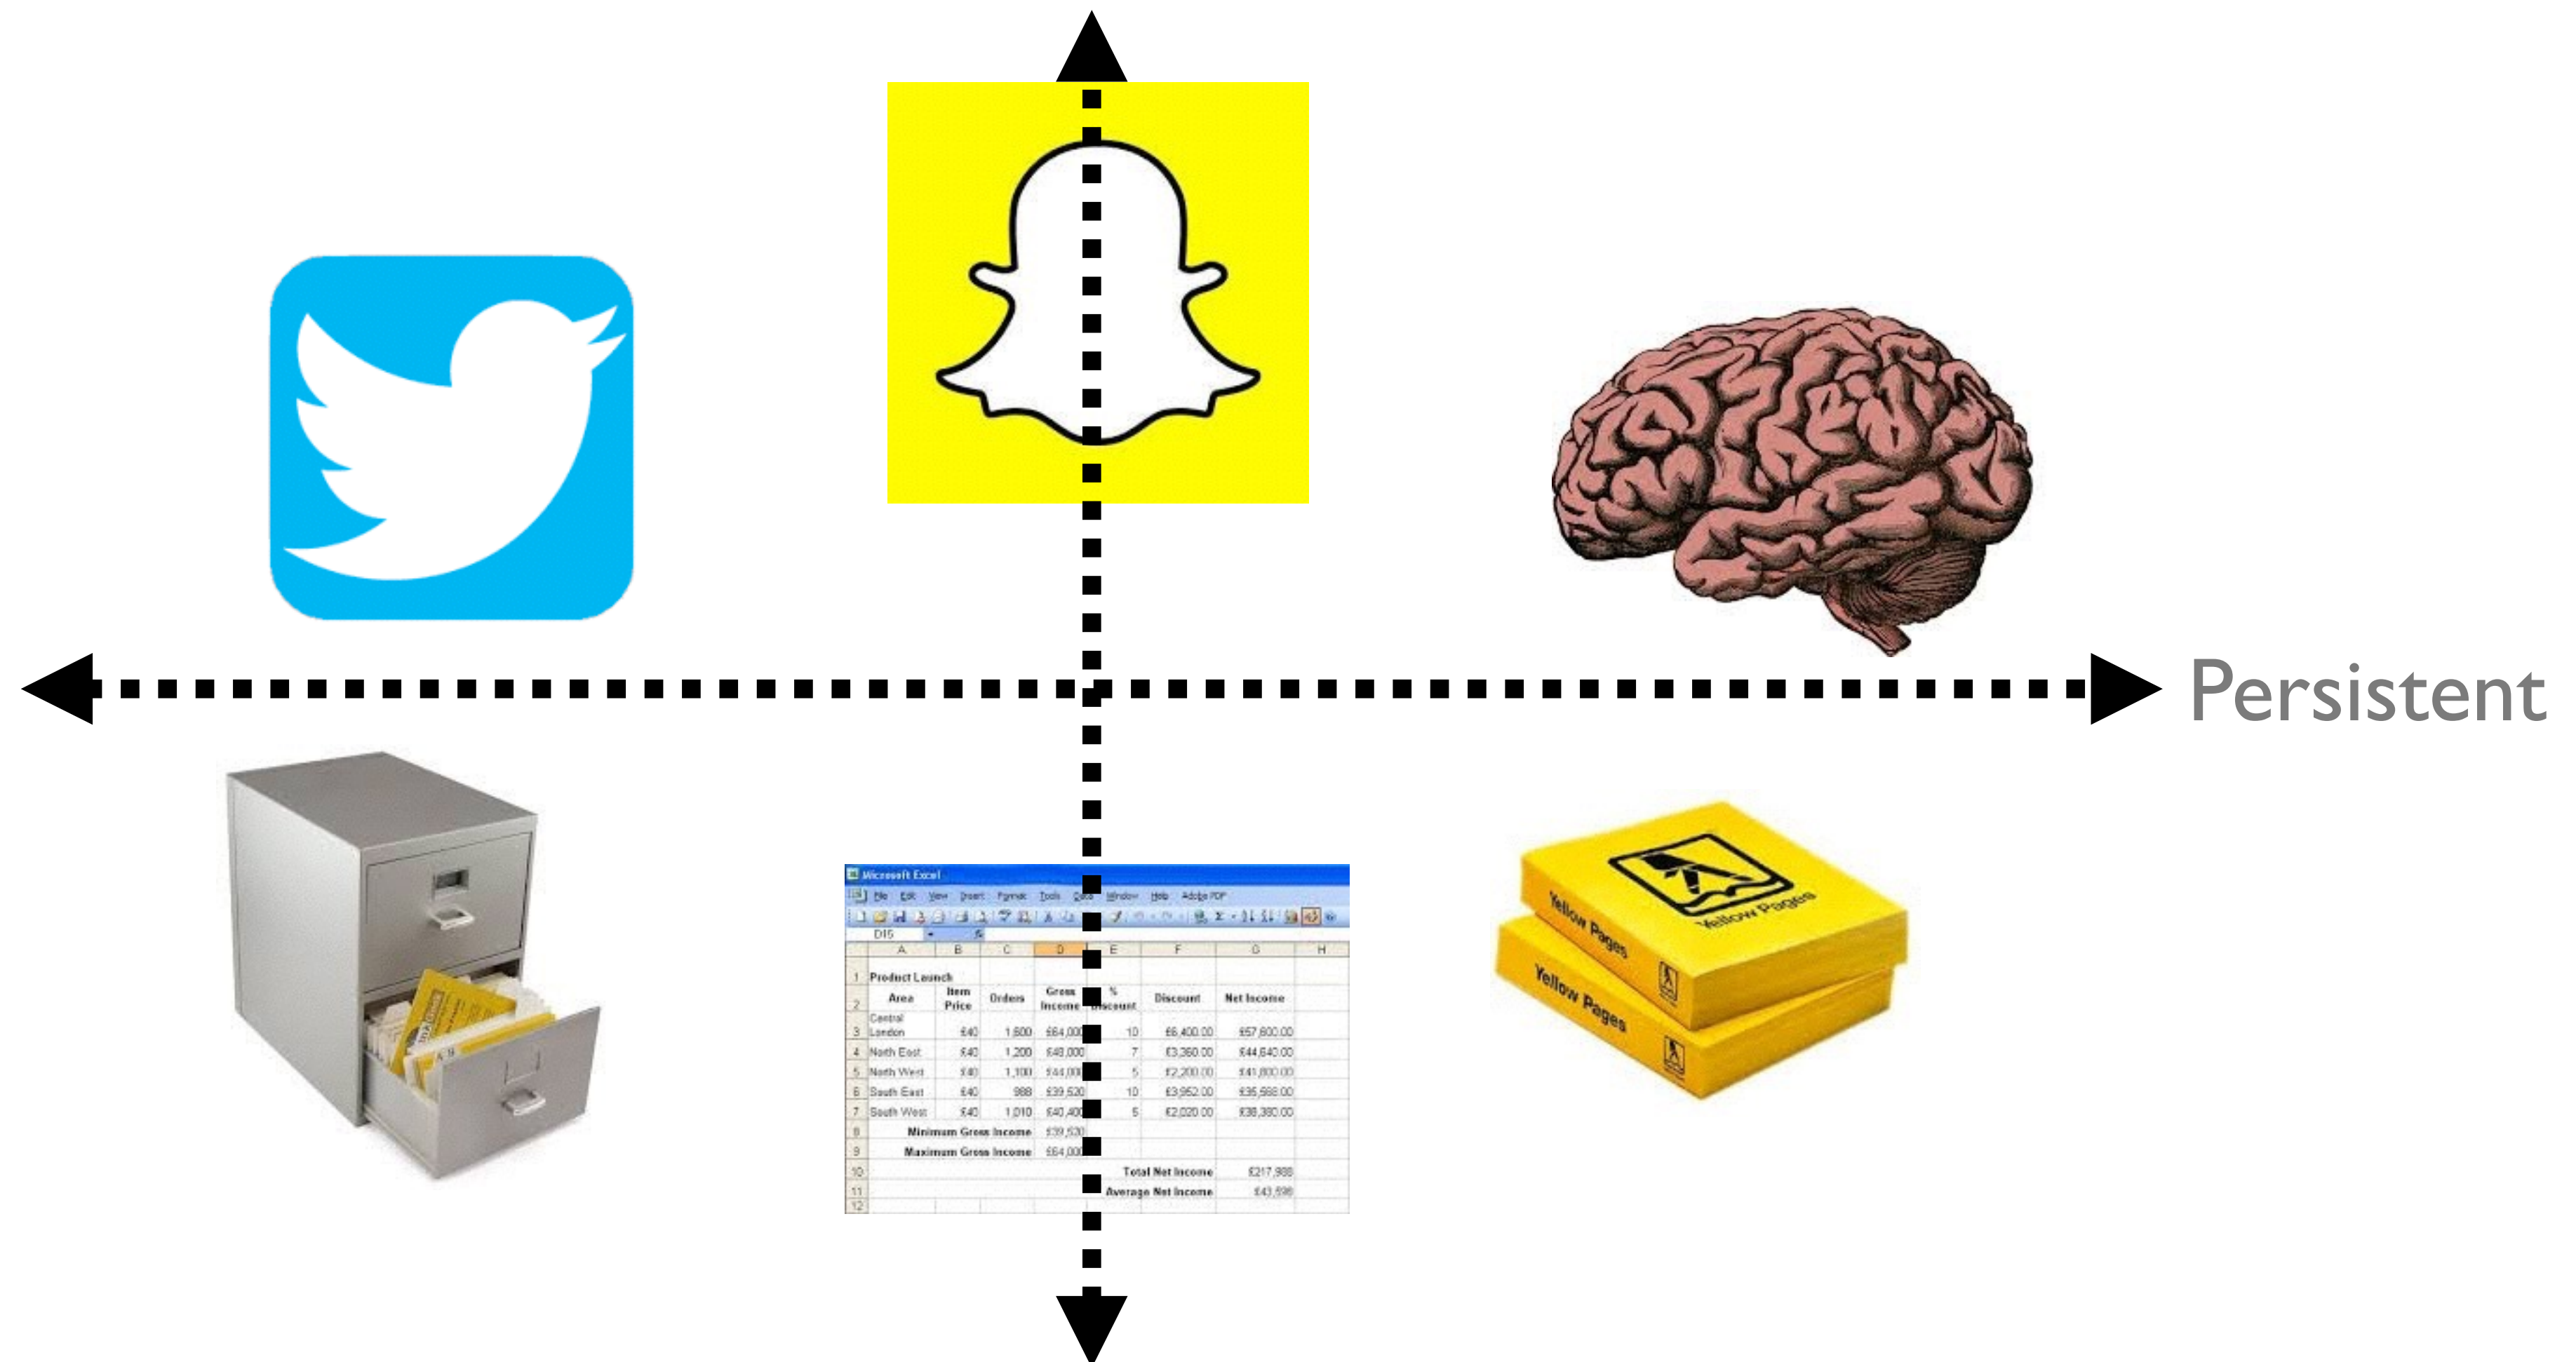
Intro to Databases

SQL

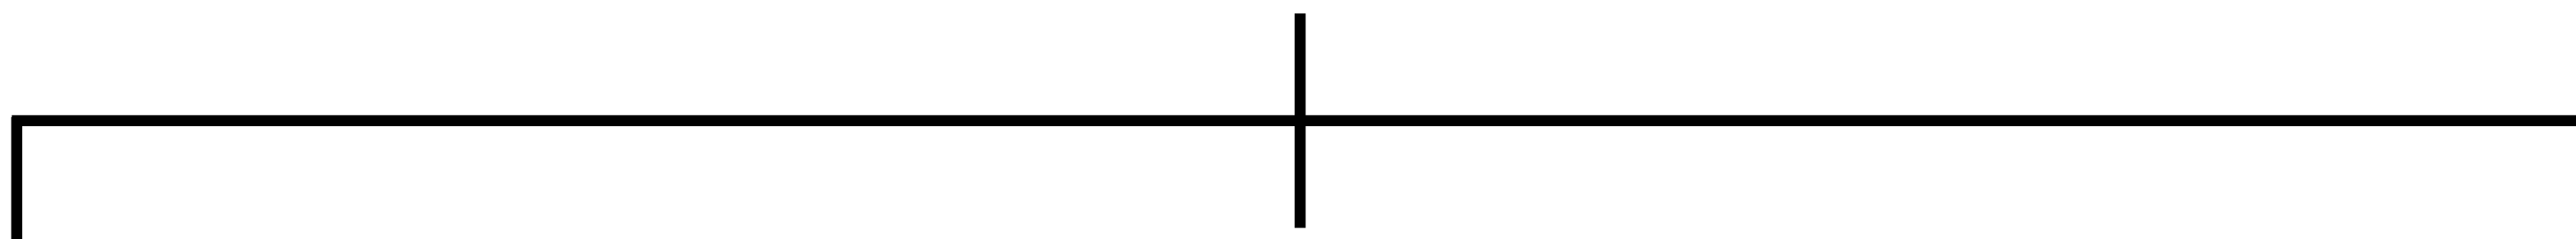
What is a database?

Things that hold info

Accessible



A database **persists** information
and is **accessible** via code



organized queryable manageable

Organized: Standard Storage Formatting

- Relational **DBs** are a collection of **Tables** (or *relations*)
- Tables have **Columns** (*attributes* / *fields*) that describe **Rows** (*instances* / *tuples*)
- Duplicate rows are not allowed
- Rows often have a **primary key** (unique identifier)

Table / Relation

Column / Attribute / Field Column / Attribute / Field Column / Attribute / Field

	ID	Name	Type
Row / Tuple / Instance	1	Pikachu	lightning
Row / Tuple / Instance	2	Squirtle	water
Row / Tuple / Instance	3	Charmander	fire
Row / Tuple / Instance	4	Bulbasaur	grass

Queryable: via a Standard Language

- A simple, structured query language: SQL
- Declarative (vs. imperative)
- No more hand-rolled algorithms / data structures
- DBMS picks an efficient execution strategy based on indexes, data, workload etc.



SQL

```
-- Pikachu, I choose you!  
SELECT id, name  
  FROM pokemon  
 WHERE type = 'lightning'  
 LIMIT 1;
```



Manageable: Easy, Safe, Performant

- Offloads work and requisite understanding of programming
- Knowledge is portable
- Abstraction
- Transfer data between systems
- DBMS (Database Management System) can make certain guarantees
 - prevent unsafe operations
 - built-in redundancies
 - handle multiple users, threads

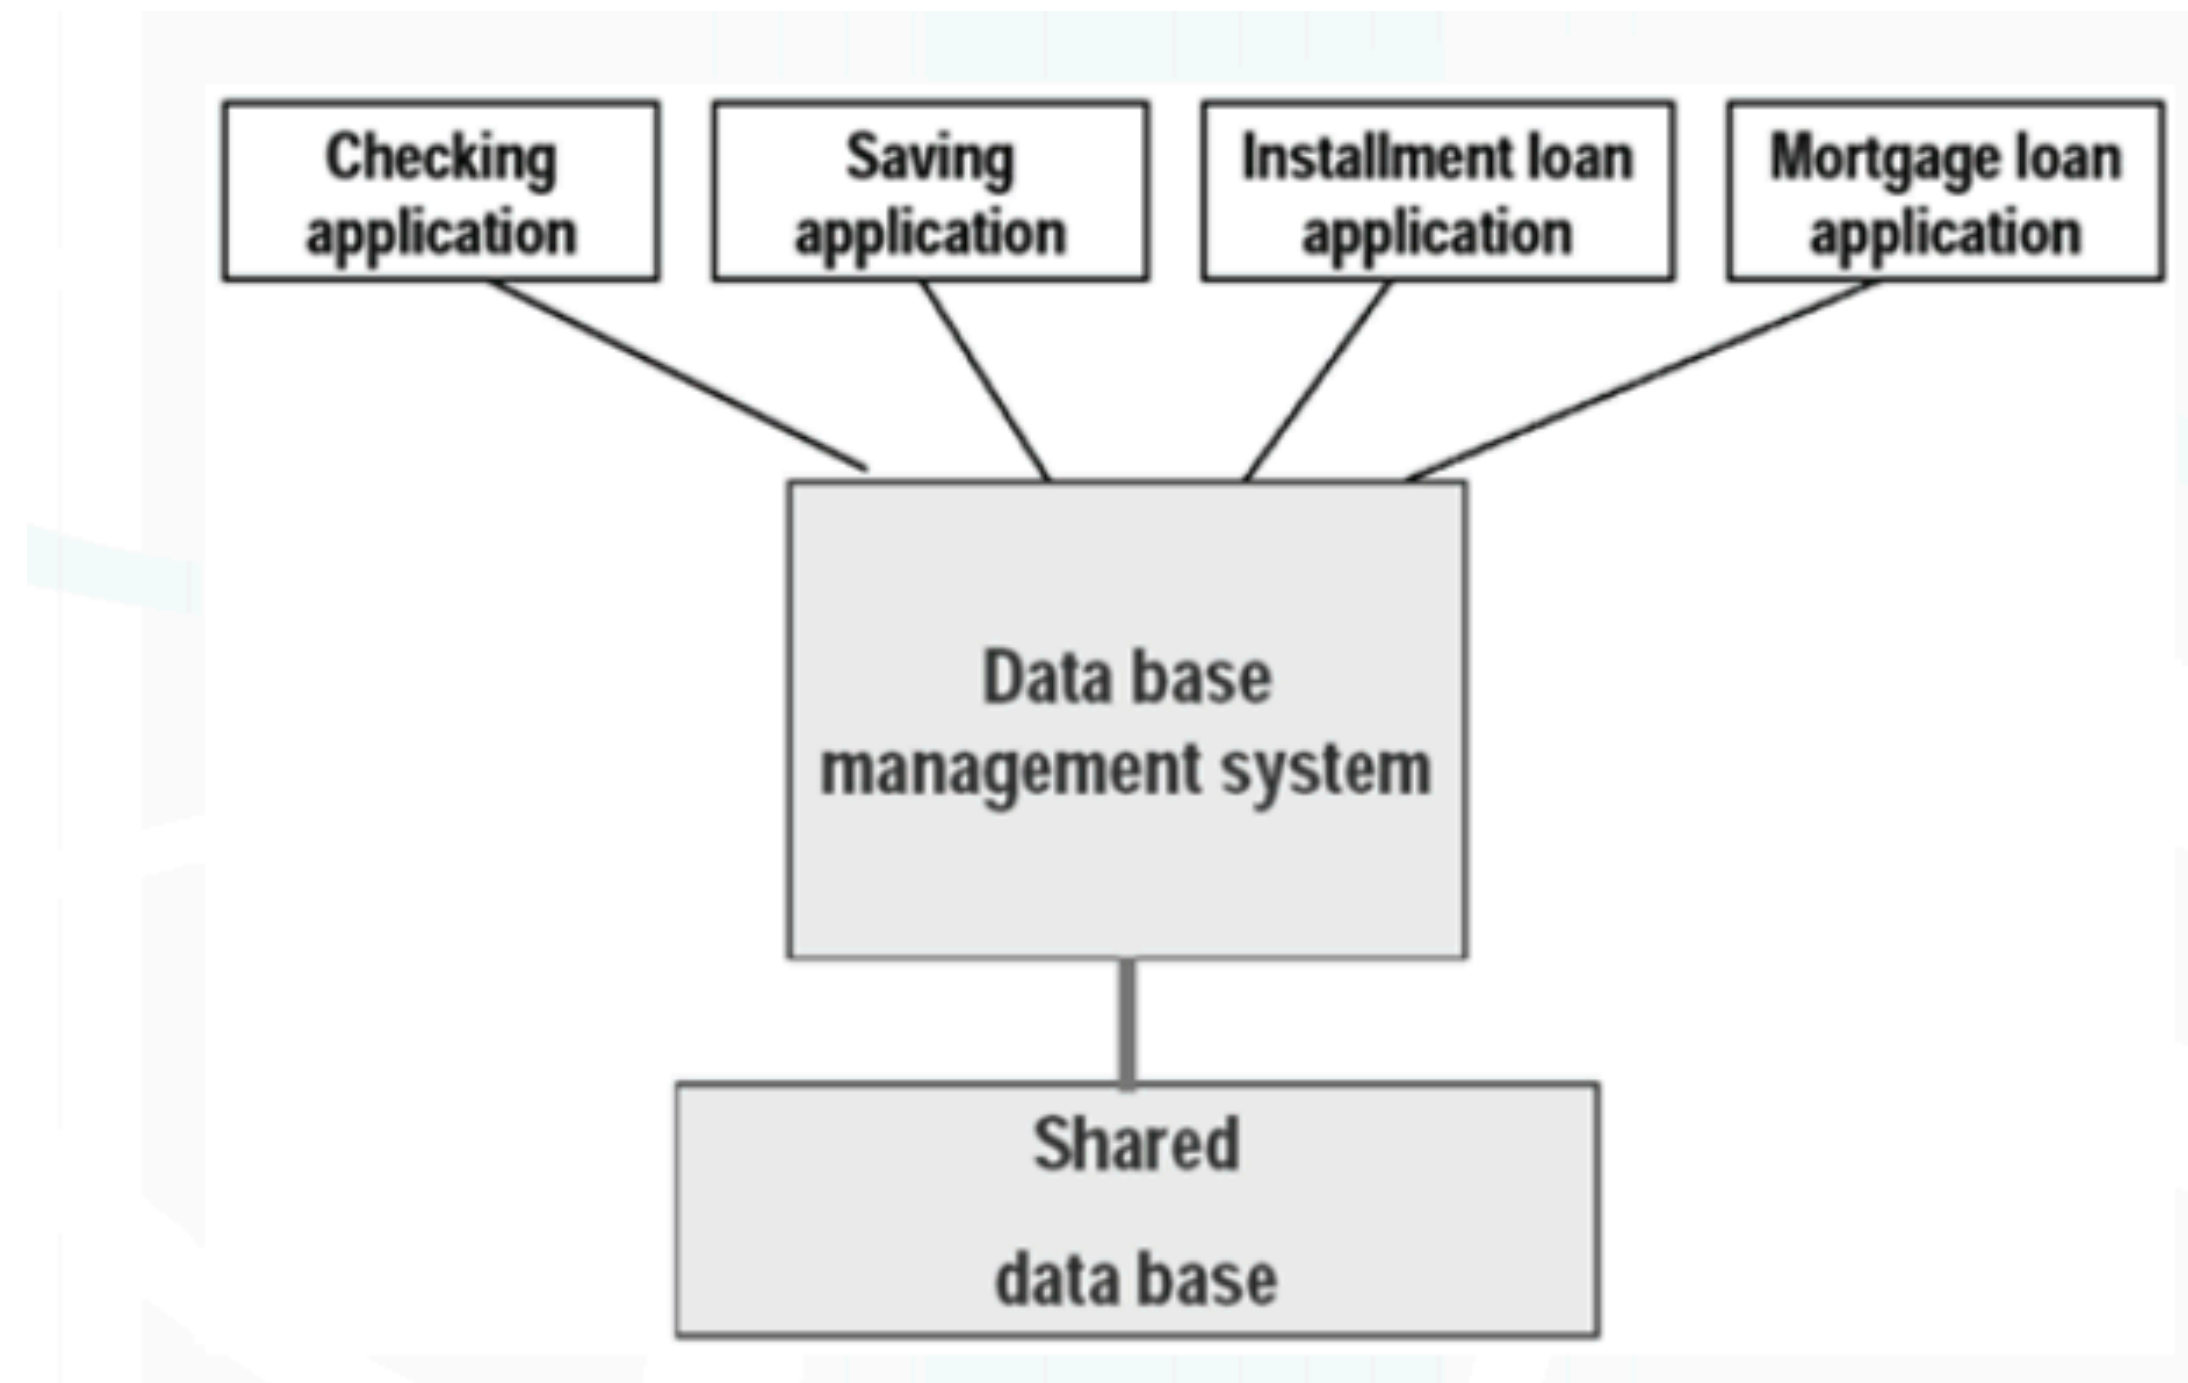
How did we end up here?

Before Relational DBs (ca. < 1970s)

- Data stored in custom “data files”
- Queried via application-specific code
- Advantages
 - Middle layer not needed
 - Solutions customized for each application
- Disadvantages
 - Hard to change the system
 - Data-transfer is difficult
 - Knowledge not compounding




Database Management Systems (DBMS)



- One layer and language to store and access data
- Sold as a way for “non-technical people” to manage data

Relational Databases & Logic



- 1969: Edgar Frank "Ted" Codd outlines *relational model* of data
- Wrote Alpha (never implemented) as a *query language*
- IBM slow to adopt his ideas
 - Competitors started to do so
 - IBM team formed without Codd, created Structured English Query Lang
- SEQUEL way better than what came before
 - 1979: copied by Larry Ellison (from pre-launch papers / talks!) as "SQL"
- SQL became the standard (ANSI 1986, ISO 1987)
 - Codd continued to fault SQL compared to his theoretical model
 - The Third Manifesto: solve the *object-relational impedance mismatch*

RDBMS vs NoSQL

- **A DBMS doesn't *have* to be relational**
 - Remember, DBMS is just an application that intelligently stores data and can answer requests to manage that data
- **Lately, many "NoSQL" or non-relational DBMSs have been gaining popularity**
 - Graph databases (e.g. Neo4j)
 - Document databases (e.g. MongoDB)
 - Hybrids (e.g. PostgreSQL)
- **RDBMSs still remain the #1 DB option for now**

Pop Quiz

1. If we didn't have “databases” what could/would data persistence look like?
2. What is a “relation”? What does it mean if a database is non-relational?
3. What is SQL and why is it important?

Appreciating Databases

- Ubiquitous
- Standardized
- Complex / deep
- Powerful: database admins are...
 - ...feared by developers.
 - ...but also taken for granted until things break. 🙄
 - ...befriended by those who want to mine (business people & governments!)

ACID Guarantees

- **Atomicity**
- **Consistency**
- **Isolation**
- **Durability**

What happens if the database crashes?

- Suppose we have a banking system with client accounts
- Every debit must have a matching credit
- Imagine a crash results in only one table being updated
 - Database inconsistency → unexpected data and software crashes/bugs
 - Financial risk for clients

Transaction Code Example

```
BEGIN;  
UPDATE accounts SET balance = balance - 100.00  
    WHERE name = 'Alice';  
UPDATE accounts SET balance = balance + 100.00  
    WHERE name = 'Bob';  
COMMIT;
```

<https://www.postgresql.org/docs/8.3/static/tutorial-transactions.html>

Atomic Transactions

- **atomic transaction: A set of database operations that must occur together**
 - i.e. A debit to one bank account, and a credit to another
- **A transaction must either succeed or fail; it cannot partially complete.**
- **Every database query is represented by a transaction**

Consistency

- **Specify rules that columns need to follow**
 - Gender column can only contain M, F, or U.
 - Savings account must start with S or checking with C
 - Column cannot be null
- **Protect the database from inconsistencies and simplify software logic**
 - Allows software to make assumptions about underlying data

Transaction Code Example

```
BEGIN;  
UPDATE accounts SET balance = balance - 100.00  
    WHERE name = 'Alice';  
UPDATE accounts SET balance = balance + 100.00  
    WHERE name = 'Bob';  
COMMIT;
```

<https://www.postgresql.org/docs/8.3/static/tutorial-transactions.html>

Resource Management

- Processes can be readers and writers
- Files can have many readers
- If a process has a writer, no other process can read from it, and no other process can write to it

Proposed File Scheme

- Suppose that we have decided not to use a database and instead store our data in a series of files.
- How might our setup fail to serve queries from multiple users?

Deadlock



Databases give us concurrency (Isolation)

- Multiple clients can make queries to read and update without the risk of deadlock or starvation.

Persistence/Durability

- Files are also persistence (store information without power)

Enough Theory. Examples!

Example DB

Primary Key
(PK)

Students

Foreign Key
(FK)

Addresses

ID	Name	Age	Gender	Address
1	Nick D.	20	M	2
2	Andy D.	28	M	2
3	Beth M.	23	F	1
4	Lisa N.	20	F	4

ID	Street	Zip	City	State
1	423 Main St.	60647	Chicago	IL
2	13 Main St	60655	Barrington	IL
3	15 Main St	60651	Elsewhere	IL
4	14 Main St	60650	Chicago	IL

All 20 Year Old Students

Students

ID	Name	Age	Gender	Address
1	Nick D.	20	M	2
2	Andy D.	28	M	2
3	Beth M.	23	F	1
4	Lisa N.	20	F	4

Result Set:

20 Year Old Students

ID	Name	Age
1	Nick D.	20
4	Lisa N.	20

```
SELECT ID, Name, Age
FROM Students
WHERE Age = 20;
```

Students

ID	Name	Age	Gender	Address
1	Nick D.	20	M	2
2	Andy D.	28	M	2
3	Beth M.	23	F	1
4	Lisa N.	20	F	4

Addresses

ID	Street	Zip	City	State
1	423 Main St.	60647	Chicago	IL
2	13 Main St.	60655	Barrington	IL
3	15 Main St.	60651	Elsewhere	IL
4	14 Main St.	60650	Chicago	IL

```
SELECT Students.ID, Name, Street, Zip, City
FROM Students
JOIN Addresses
ON Students.Address = Addresses.ID;
```

Result Set:

Students with Addresses

Student.ID	Name	Street	Zip	City
1	Nick D.	13 Main St.	60655	Barrington
2	Andy D.	13 Main St.	60655	Barrington
3	Beth M.	423 Main St.	60647	Chicago
4	Lisa N.	14 Main St.	60650	Chicago

Students

ID	Name	Age	Gender	Address
1	Nick D.	20	M	2
2	Andy D.	28	M	2
3	Beth M.	23	F	1
4	Lisa N.	20	F	4

Addresses

ID	Street	Zip	City	State
1	423 Main St.	60647	Chicago	IL
2	13 Main St.	60655	Barrington	IL
3	15 Main St.	60651	Elsewhere	IL
4	14 Main St.	60650	Chicago	IL

```
SELECT Students.ID, Name, Street, Zip, City
FROM Students
JOIN Addresses
  ON Students.Address = Addresses.ID
WHERE Addresses.City = 'Chicago';
```

Result Set:

Students with Addresses

Student.ID	Name	Street	Zip	City
3	Beth M.	423 Main St.	60647	Chicago
4	Lisa N.	14 Main St.	60650	Chicago




Some Common SQL Keywords

Keyword	Action
SELECT	Which COLUMNS to include in output table (shrinks the result horizontally!)
FROM	Which TABLE to pull data from
JOIN	Another TABLE to glue / concatenate to the output
ON	What COLUMNS must match when joining two tables
WHERE	Which ROWS to include in the output table (shrinks the result vertically!)

Schema and Content

- Schema: table's blueprint for data shape/format
- Content: actual data (a row) e.g. {1, "Bart S.", 10, "M"}
- A schema is used to validate incoming content



SELECT * FROM PEOPLE WHERE AGE > 30





QUESTION

 `SELECT * FROM PEOPLE WHERE ANNUAL_INCOME > 100 000`



 `SELECT * FROM PEOPLE WHERE AFRAID_OF_FLYING = TRUE`





SQL

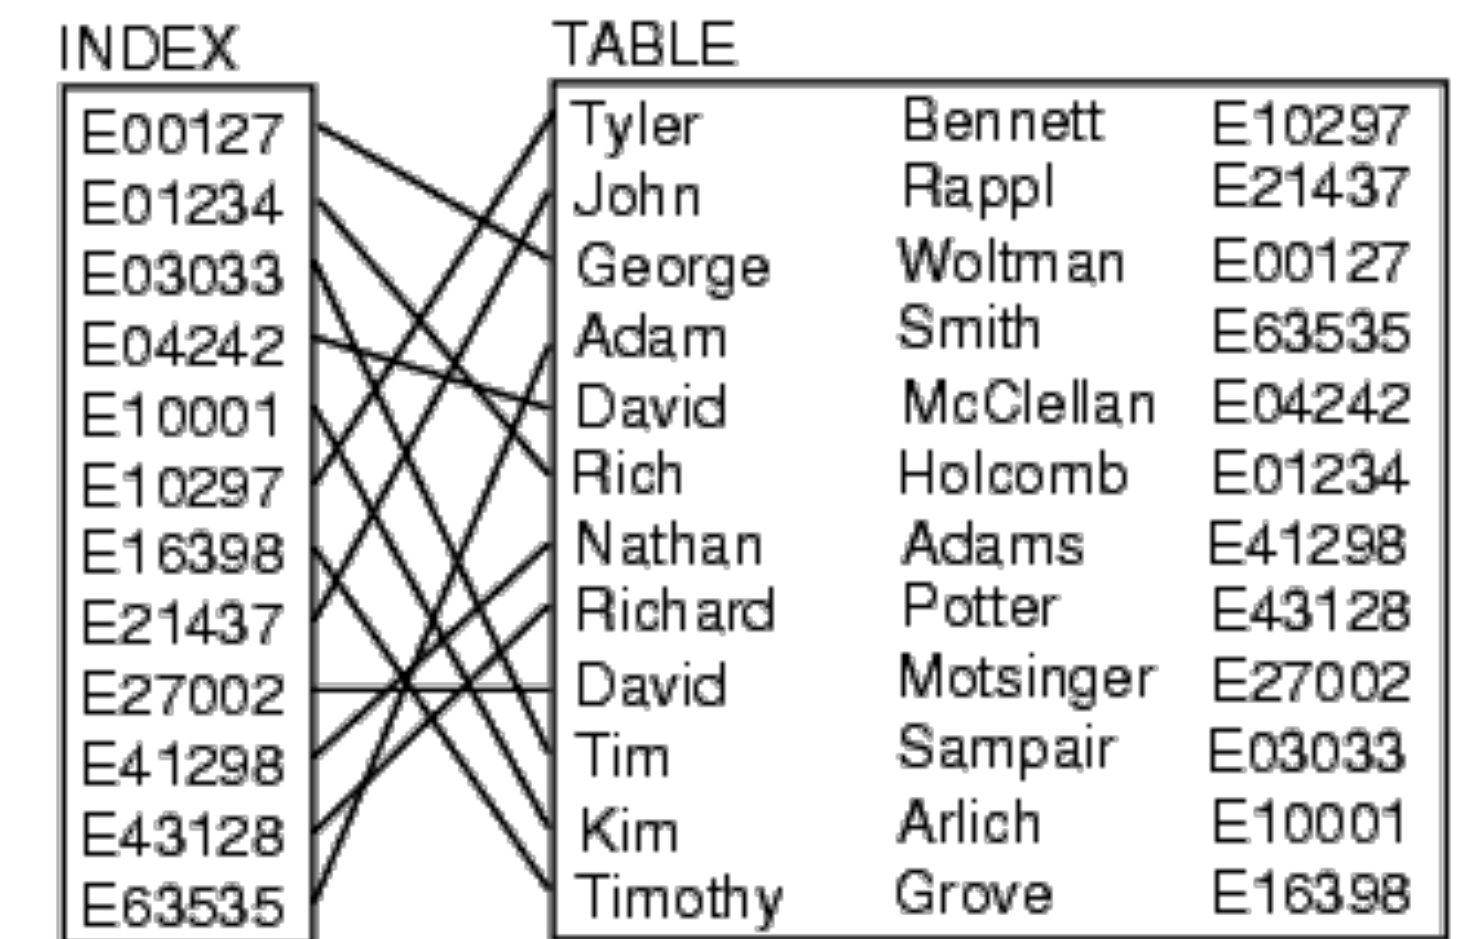
SQL is used to create/read/update/delete (CRUD) data from a database

- **INSERT:** Insert new rows into a table
- **SELECT:** Get data from a database
- **UPDATE:** Update existing rows in a table
- **DELETE:** Delete rows from a table

- **CREATE / DROP:** Make / delete new dbs/tables/views/indexes

Quick Notes on Indexing

- Creates another data structure, which holds
 - row's indexed field's value
 - a pointer to the record the row relates to
- **Index on Employee_ID**
 - Index points to record
 - Index is ordered



<https://www.progress.com/tutorials/odbc/using-indexes>

Quick Notes on Indexing

- A database driver can use indexes to find records quickly
- Without an index, the driver's worst case is searching the entire database table to find a record for a specific `Employee_ID`
- Downside
 - Extra storage space requirements (though less than a copy of the table)
 - Redundant data means update/insert/delete takes more time



<https://lol.browserling.com/tables.png>

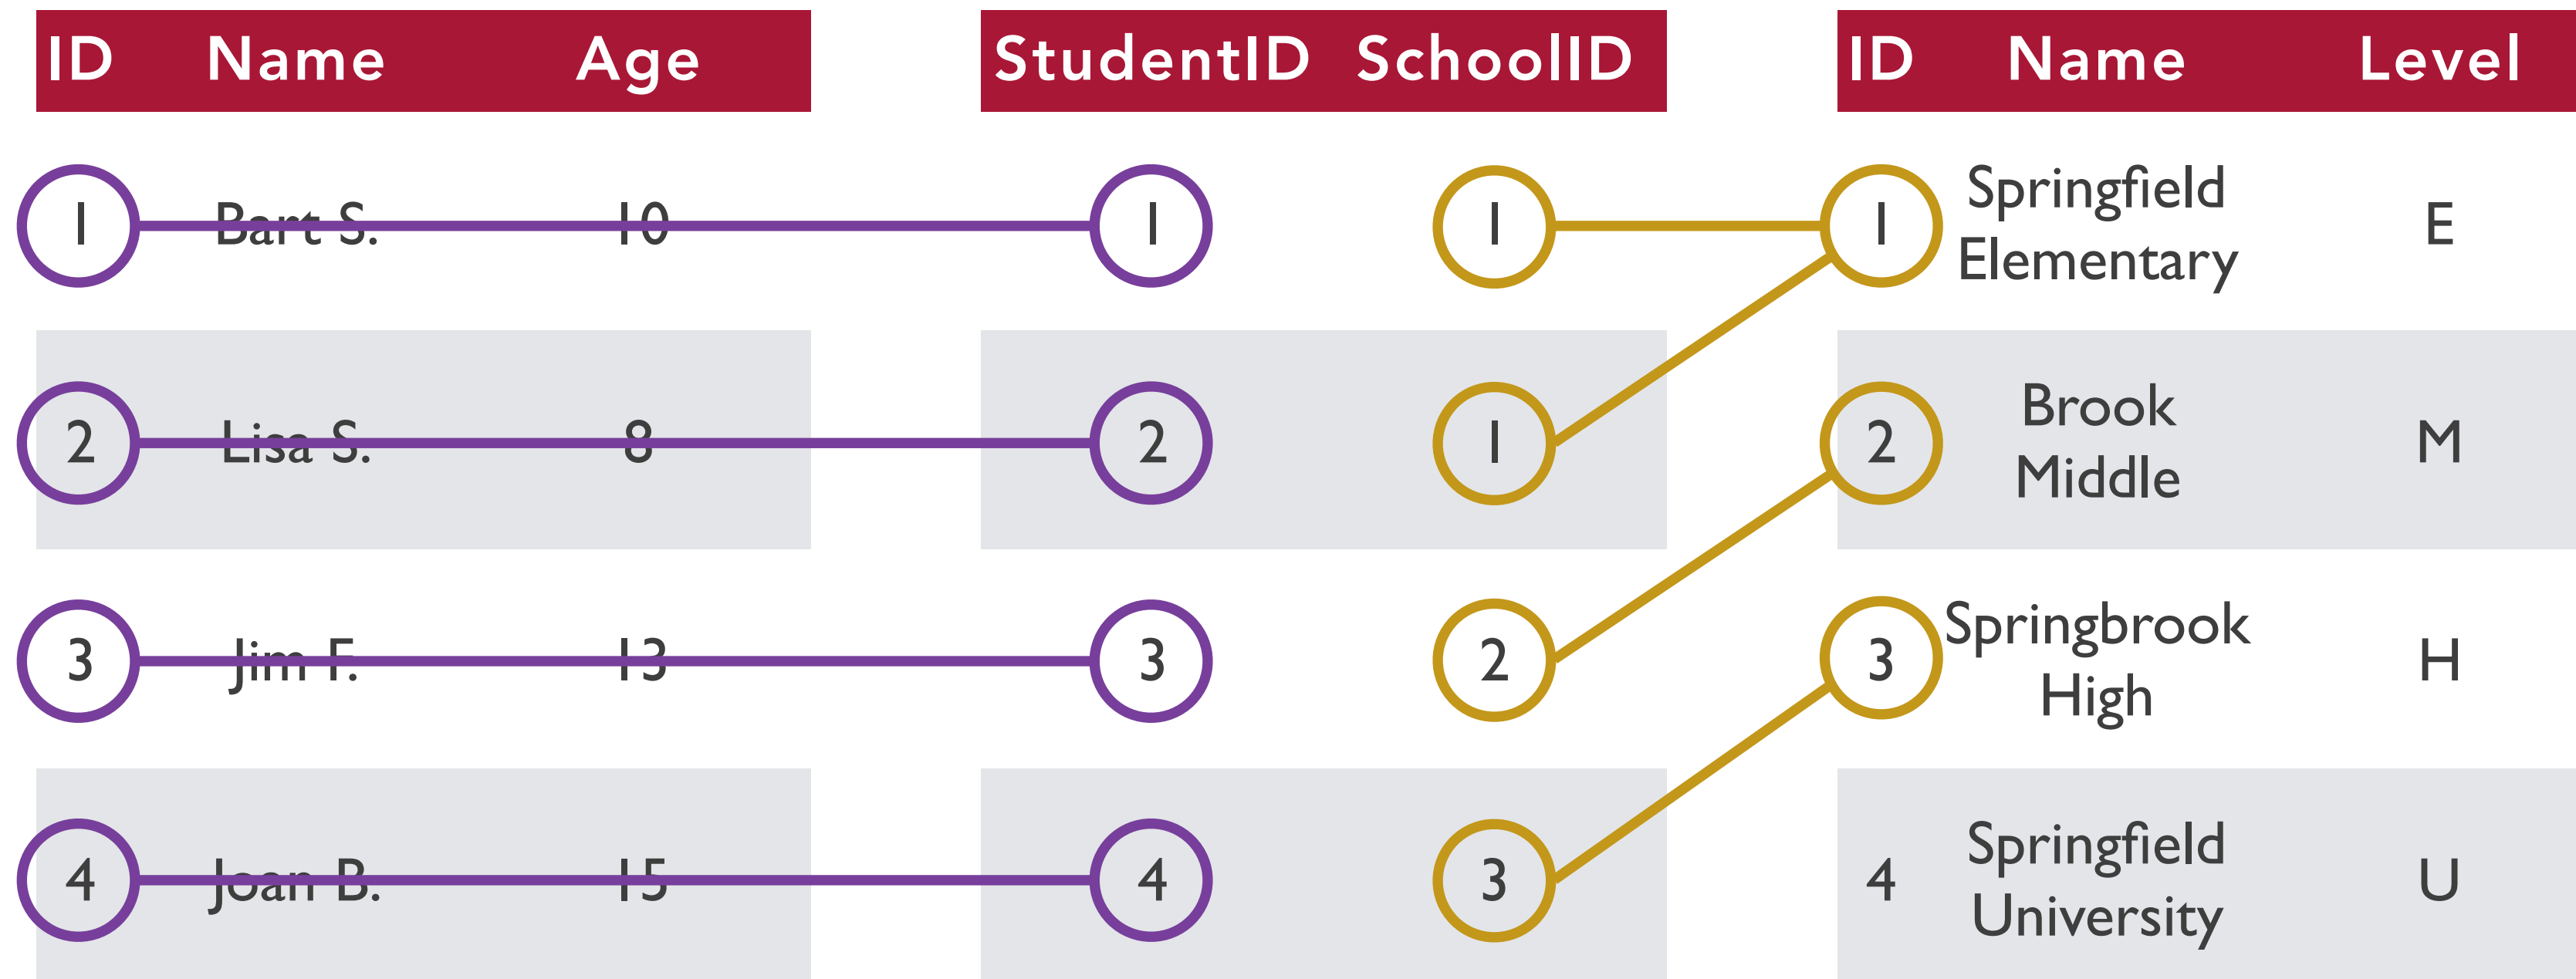
**Find all students from
Springfield Elementary**



Student

Enrollment

School



```
SELECT *  
FROM Student  
INNER JOIN Enrollment  
ON Student.ID = Enrollment.StudentID  
INNER JOIN School  
ON Enrollment.SchoolID = School.ID;
```




SELECT *

FROM Student

ID	Name	Age	StudentID	SchoolID
----	------	-----	-----------	----------

INNER JOIN Enrollment

ON Student.ID = Enrollment.StudentID

INNER JOIN School

ON Enrollment.SchoolID = School.ID;

ID	Name	Level
----	------	-------

1 Springfield Elementary E

2 Brook Middle M

Student.ID Student.Name Age StudentID SchoolID School.ID School.Name Level

1	Bart S.	10	1	1	1	Springfield Elementary	E
2	Lisa S.	8	2	1	1	Springfield Elementary	E
3	Jim F.	13	3	2	2	Brook Middle	M
4	Joan B.	15	4	3	3	Springbrook High	H




```
SELECT *
FROM Student
INNER JOIN Enrollment
    ON Student.ID = Enrollment.StudentID
INNER JOIN School
    ON Enrollment.SchoolID = School.ID;
WHERE School.Name = 'Springfield Elementary';
```

Student.ID	Student.Name	Age	StudentID	SchoolID	School.ID	School.Name	Level
1	Bart S.	10	1	1	1	Springfield Elementary	E
2	Lisa S.	8	2	1	1	Springfield Elementary	E




Inner Join




```
SELECT *  
FROM A  
INNER JOIN B  
ON A.Key = B.Key
```

Outer Join




```
SELECT *  
FROM A  
FULL OUTER JOIN B  
ON A.Key = B.Key
```

Left Join



```
SELECT *  
FROM A  
LEFT JOIN B  
ON A.Key = B.Key
```

Right Join




```
SELECT *  
FROM A  
RIGHT JOIN B  
ON A.Key = B.Key
```

<http://www.codeproject.com/Articles/33052/Visual-Representation-of-SQL-Joins>



Inner Join



```
SELECT pets.name, owners.name
FROM owners
INNER JOIN pets
ON pets.ownerID = owners.ID
```

OWNERS

ID	name
1	Geordi
2	Janeway
3	Data
4	Spok

PETS

ID	ownerID	type	name
1	4	Monkey	Mittens
2	null	Lizard	Carol
3	1	Dog	Rufus
4	2	Cat	Fireball

pets.name	owners.name
Mittens	Spok
Rufus	Geordi
Fireball	Janeway




PETS

ID	ownerID	type	name
1	4	Monkey	Mittens
2	null	Lizard	Carol
3	1	Dog	Rufus
4	2	Cat	Fireball

pets.name	owners.name
Mittens	Spok
Rufus	Geordi
Fireball	Janeway
null	Data




Left Join



```
SELECT pets.name, owners.name  
FROM owners  
LEFT JOIN pets  
ON pets.ownerID = owners.ID
```

OWNERS

ID	name
1	Geordi
2	Janeway
3	Data
4	Spok



pets.name	owners.name
Mittens	Spok
Carol	null
Rufus	Geordi
Fireball	Janeway


OWNERS

ID	name
1	Geordi
2	Janeway
3	Data
4	Spok

PETS

ID	ownerID	type	name
1	4	Monkey	Mittens
2	null	Lizard	Carol
3	1	Dog	Rufus
4	2	Cat	Fireball

Right Join




```
SELECT pets.name, owners.name  
FROM owners  
RIGHT JOIN pets  
ON pets.ownerID = owners.ID
```



OWNERS

ID	name
1	Geordi
2	Janeway
3	Data
4	Spok

Outer Join



```
SELECT pets.name, owners.name  
FROM owners  
FULL OUTER JOIN pets  
ON pets.ownerID = owners.ID
```

PETS

	pets.name	owners.name
	Mittens	Spok
➔	Carol	null
	Rufus	Geordi
	Fireball	Janeway
➔	null	Data

ID	ownerID	type	name
1	4	Monkey	Mittens
2	null	Lizard	Carol
3	1	Dog	Rufus
4	2	Cat	Fireball

Pop Quiz

1. What is ACID compliance and why is it important?
2. Give examples of how to do all CRUD operations in SQL.
3. Name the different types of joins and describe the differences between them.

SQL POWER



AS

ID	Name	Age
1	Bart S.	10
2	Lisa S.	8
3	Jim F.	13
4	Joan B.	15

StudentID	SchoolID
1	1
2	1
3	2
4	3

ID	Name	Level
1	Springfield Elementary	E
2	Brook Middle	M
3	Springbrook High	H
4	Springfield University	U

```
SELECT *
FROM Student AS st
INNER JOIN Enrollment AS e
ON st.ID = e.StudentID
INNER JOIN School as sc
ON e.SchoolID = sc.ID;
```

st.ID	st.Name	Age	StudentID	SchoolID	sc.ID	sc.Name	Level
1	Bart S.	10	1	1	1	Springfield Elementary	E
2	Lisa S.	8	2	1	1	Springfield Elementary	E
3	Jim F.	13	3	2	2	Brook Middle	M
4	Joan B.	15	4	3	3	Springbrook High	H



ID	Name	Age
1	Bart S.	10
2	Lisa S.	8
3	Jim F.	13
4	Joan B.	15

StudentID	SchoolID
1	1
2	1
3	2
4	3

ID	Name	Level
1	Springfield Elementary	E
2	Brook Middle	M
3	Springbrook High	H
4	Springfield University	U

AS (without AS)

```
SELECT *
FROM Student st
INNER JOIN Enrollment e
ON st.ID = e.StudentID
INNER JOIN School sc
ON e.SchoolID = sc.ID;
```

st.ID	st.Name	Age	StudentID	SchoolID	sc.ID	sc.Name	Level
1	Bart S.	10	1	1	1	Springfield Elementary	E
2	Lisa S.	8	2	1	1	Springfield Elementary	E
3	Jim F.	13	3	2	2	Brook Middle	M
4	Joan B.	15	4	3	3	Springbrook High	H



GROUP BY

+

COUNT

ID	Name	Age
1	Bart S.	10
2	Lisa S.	8
3	Jim F.	13
4	Joan B.	15

StudentID	SchoolID
1	1
2	1
3	2
4	3

ID	Name	Level
1	Springfield Elementary	E
2	Brook Middle	M
3	Springbrook High	H
4	Springfield University	U

```
SELECT Name, COUNT(*)
FROM School
INNER JOIN Enrollment
ON School.ID = Enrollment.StudentID
GROUP BY Name;
```

Name	COUNT(*)
Springfield Elementary	2
Brook Middle	1
Springbrook High	1
Springfield University	0



ORDER BY

ID	Name	Age
1	Bart S.	10
2	Lisa S.	8
3	Jim F.	13
4	Joan B.	15

StudentID	SchoolID
1	1
2	1
3	2
4	3

ID	Name	Level
1	Springfield Elementary	E
2	Brook Middle	M
3	Springbrook High	H
4	Springfield University	U

```
SELECT *
FROM Student
ORDER BY Age DESC;
```

ID	Name	Age
4	Joan B.	15
3	Jim F.	13
1	Bart S.	10
2	Lisa S.	8



SUB-QUERIES

```
SELECT ID, Name, Age
FROM Student
INNER JOIN Enrollment
  ON Student.ID = Enrollment.StudentID
INNER JOIN (
  SELECT SchoolID
  FROM Student
  WHERE Student.Name = 'Lisa S.'
  INNER JOIN Enrollment
    ON Student.ID = Enrollment.StudentID
) AS LisaSchools
  ON LisaSchools.SchoolID = Enrollment.SchoolID
WHERE Name != 'Lisa S.';
```

ID	Name	Age
1	Bart S.	10
2	Lisa S.	8
3	Jim F.	13
4	Joan B.	15

StudentID	SchoolID
1	1
2	1
3	2
4	3

ID	Name	Level
1	Springfield Elementary	E
2	Brook Middle	M
3	Springbrook High	H
4	Springfield University	U

ID	Name	Age
1	Bart S.	10

WORKSHOP